

Project 5

December 2, 2019

```
[7]: import numpy as np
import pandas as pd
from IPython.display import display, Latex

#Define our target function, general regression equation, and arrays for
→building the data sets#
def target(x):
    return np.square(x) + 10

def sum_square_error(w,y):
    return np.square(w-y)

j = 12
test_j = 5
x_train = np.zeros(shape = (j,3))
y_train = np.zeros(shape = (j,1))

#Developing the training set#
np.random.seed(123)
get_uniform_values = np.random.uniform(-2,10,j)

for row in range(0,j):
    x_train[row] = (1,get_uniform_values[row],np.square(get_uniform_values[row]))

y_train = target(x_train[:,1])

#Developing the test set and ensuring it shares no points with the training set#
x_test = np.zeros(shape = (test_j,3))
get_uniform_values = np.random.uniform(-2,10,test_j)

for row in range(0,test_j):
    x_test[row] = (1,get_uniform_values[row],np.square(get_uniform_values[row]))
    check_for_common_values = np.isin(x_test,x_train[:,1])
    while check_for_common_values is True:
        get_uniform_values = np.random.uniform(-2,10,test_j)
        x_test[row] = (1,get_uniform_values[row],np.
→square(get_uniform_values[row]))
```

```

y_test = target(x_test[:,1])

x_train_display = np.column_stack((x_train[:,1],y_train))
train_df = pd.DataFrame(x_train_display,columns = ['x','y'])
print("Part a. \n\nTraining Data Set\n", train_df)

x_test_display = np.column_stack((x_test[:,1],y_test))
test_df = pd.DataFrame(x_test_display,columns = ['x','y'])
print("\n\nTest Data Set\n", test_df)

```

Part a.

Training Data Set

	x	y
0	6.357630	50.419462
1	1.433672	12.055415
2	0.722217	10.521598
3	4.615777	31.305399
4	6.633628	54.005016
5	3.077278	19.469637
6	9.769170	105.436690
7	6.217957	48.662988
8	3.771183	24.221820
9	2.705410	17.319244
10	2.118136	14.486501
11	6.748596	55.543555

Test Data Set

	x	y
0	3.262867	20.646301
1	-1.283865	11.648310
2	2.776531	17.709125
3	6.855945	57.003980
4	0.189901	10.036062

```

[8]: #Get the weights for our quadratic regression#
x_train_inv = np.linalg.pinv(x_train)
weights = np.matmul(x_train_inv,y_train)

#Compute the errors over our data sets#
weighted_average_train = weighted_average_test = error_train = error_test = 0

for row in range(0,j):
    weighted_average_train = np.dot(weights,x_train[row])
    error_train += sum_square_error(weighted_average_train,y_train[row])

```

```

for row in range(0,test_j):
    weighted_average_test = np.dot(weights,x_test[row])
    error_test += sum_square_error(weighted_average_test,y_test[row])

error_train *= (1/j)
error_test *= (1/test_j)
print("Part b. \n Quadratic Regression Equation is roughly: " )
display(Latex("$\hat{y}=10x_0+0x_1+1x_1^2$"))
print("Training Error for Quadratic Regression: ", error_train)
print("Test Error for Quadratic Regression: ", error_test)

```

Part b.

Quadratic Regression Equation is roughly:

$$\hat{y} = 10x_0 + 0x_1 + 1x_1^2$$

Training Error for Quadratic Regression: 9.090307164496792e-28

Test Error for Quadratic Regression: 5.086575116865085e-28

```

[9]: #####Compute the linear regresssion using ridge regression#####
def ridge_penalty(lam,w):
    return lam * np.sum(np.square(w))

#Bring x_train back down to linear terms and construct lambda to hold the
→penalty#
#and its associated cross validation error#
x_train = np.delete(x_train,2,1)
lam_ridge = np.array([[0.1,0],[1,0],[10,0],[100,0]])
lam_lasso = np.array([[0.1,0],[1,0],[10,0],[100,0]])

#Test each lambda on each of the cross_validation sets#
for penalty in range(0,len(lam_ridge)):
    error_cross_validate = error_train_ridge = error_test_ridge = 0
#Construct each cross validation set to test lambda on#
    k_fold = 0
    while k_fold <= 8:
        x_cross_validate = np.delete(x_train,slice(k_fold,k_fold+4),0)
        y_cross_validate = np.delete(y_train,slice(k_fold,k_fold+4),0)
        x_train_inv = np.linalg.pinv(x_cross_validate)
        weights = np.matmul(x_train_inv,y_cross_validate)

#Calculate the training error on each validation set#
        for row in range(0,len(x_cross_validate)):
            weighted_average_train = np.dot(weights,x_cross_validate[row])
            error_train_ridge += sum_square_error(weighted_average_train,
→y_cross_validate[row])

```

```

#Calculate the mean square error of the cross validation set#
    error_train_ridge *= (1/len(x_cross_validate))
#Add the penalty to our training error#
    error_train_ridge += ridge_penalty(lam_ridge[penalty][0],weights)
#Add the new training error to our cross validation error#
    error_cross_validate += error_train_ridge

    k_fold += 4
#Take the average of the training errors and record as cross validation error#
    error_cross_validate *= (1/3)
#Store the cross validation error in the lambda array from above#
    lam_ridge[penalty][1] = error_cross_validate

#Determine which lambda gave us the best cross validation error#
error_cross_validate_ridge = np.delete(lam_ridge,0,1)
lambda_optimal = np.argmin(error_cross_validate_ridge)

#Validate our chosen model on the entire training set#
x_train_inv = np.linalg.pinv(x_train)
weights = np.matmul(x_train_inv,y_train)

for row in range(0,j):
    weighted_average_train = np.dot(weights,x_train[row])
    error_train_ridge += sum_square_error(weighted_average_train, y_train[row])

error_train_ridge *= (1/j)
error_train_ridge += ridge_penalty(lam_ridge[lambda_optimal][0],weights)
lambda_ridge_df = pd.DataFrame(lam_ridge,columns = ['lambda','Cross-Validation_
    ↳Error'])
print("Part c. \nCross-Validation Errors obtained during ridge_
    ↳regression\n\n",lambda_ridge_df, "\n\n Cross-Validation Errors obtained during_
    ↳lasso regression\n")

#Calculate the error on the test set#
x_test = np.delete(x_test,2,1)
for row in range(0,test_j):
    weighted_average_test = np.dot(weights,x_test[row])
    error_test_ridge += sum_square_error(weighted_average_test,y_test[row])

error_test_ridge *= (1/test_j)
error_test_ridge += ridge_penalty(lam_ridge[lambda_optimal][0],weights)

#####Compute Lasso Regression#####
def lasso_penalty(lam,w):
    return lam * np.sum(np.abs(w))

for penalty in range(0,len(lam_lasso)):

```

```

    error_cross_validate = error_train_lasso = error_test_lasso = 0
#Construct each cross validation set to test lambda on#
    k_fold = 0
    while k_fold <= 8:
        x_cross_validate = np.delete(x_train,slice(k_fold,k_fold+4),0)
        y_cross_validate = np.delete(y_train,slice(k_fold,k_fold+4),0)
        x_train_inv = np.linalg.pinv(x_cross_validate)
        weights = np.matmul(x_train_inv,y_cross_validate)

#Calculate the training error on each validation set#
        for row in range(0,len(x_cross_validate)):
            weighted_average_train = np.dot(weights,x_cross_validate[row])
            error_train_lasso += sum_square_error(weighted_average_train,
            ↪y_cross_validate[row])

#Calculate the mean square error of the cross validation set#
            error_train_lasso *= (1/len(x_cross_validate))
#Add the penalty to our training error#
            error_train_lasso += lasso_penalty(lam_lasso[penalty][0],weights)
#Add the new training error to our cross validation error#
            error_cross_validate += error_train_lasso

        k_fold += 4
#Take the average of the training errors and record as cross validation error#
        error_cross_validate *= (1/3)
#Store the cross validation error in the lambda array from above#
        lam_lasso[penalty][1] = error_cross_validate

#Determine which lambda gave us the best cross validation error#
    error_cross_validate_ridge = np.delete(lam_lasso,0,1)
    lambda_optimal = np.argmin(error_cross_validate_ridge)

    del error_cross_validate_ridge, penalty, k_fold, row

#Validate our chosen model on the entire training set#
    x_train_inv = np.linalg.pinv(x_train)
    weights = np.matmul(x_train_inv,y_train)

    for row in range(0,j):
        weighted_average_train = np.dot(weights,x_train[row])
        error_train_lasso += sum_square_error(weighted_average_train, y_train[row])

    error_train_lasso *= (1/j)
    error_train_lasso += lasso_penalty(lam_lasso[lambda_optimal][0],weights)
    lambda_lasso_df = pd.DataFrame(lam_lasso,columns = ['lambda','Cross-Validation_
    ↪Error'])
    print(lambda_lasso_df)

```

```

#Calculate the error on the test set#
for row in range(0,test_j):
    weighted_average_test = np.dot(weights,x_test[row])
    error_test_lasso += sum_square_error(weighted_average_test,y_test[row])

error_test_lasso *= (1/test_j)
error_test_lasso += lasso_penalty(lam_lasso[lambda_optimal][0],weights)

print("\nPart d. \nI chose this Cross-Validation Error for Ridge Regression: ",
      lam_lasso[lambda_optimal][1])
print("I chose this Cross-Validation Error for Lasso Regression: ",
      lam_lasso[lambda_optimal][1])

```

Part c.

Cross-Validation Errors obtained during ridge regression

	lambda	Cross-Validation Error
0	0.1	60.763396
1	1.0	257.761536
2	10.0	2227.742939
3	100.0	21927.556972

Cross-Validation Errors obtained during lasso regression

	lambda	Cross-Validation Error
0	0.1	40.779563
1	1.0	57.923206
2	10.0	229.359641
3	100.0	1943.723994

Part d.

I chose this Cross-Validation Error for Ridge Regression: 60.76339570283858

I chose this Cross-Validation Error for Lasso Regression: 40.77956272511449

```

[10]: print("Part e.\n\nRidge Regression\nLinear Ridge Regression Equation is roughly:
      ↪")
display(Latex("$$E_{ridge}=\sum_{j=1}^n(-7.818x_{0j}+9.918x_{1j}-y_j)^2+0.1((-7.
      ↪818)^2+9.918^2)$$"))
print("\nTraining Error for Ridge Regression: ", error_train_lasso)
print("Test Error for Ridge Regression", error_test_lasso)
print("\n\nLasso Regression\nLinear Lasso Regression Equation is roughly: ")
display(Latex("$$E_{lasso}=\sum_{j=1}^n(-7.818x_{0j}+9.918x_{1j}-y_j)^2+0.1(|-7.
      ↪818|+|9.918|)$$"))
print("Training Error for Lasso Regression: ", error_train_lasso)
print("Test Error for Lasso Regression", error_test_lasso)

```

Part e.

Ridge Regression

Linear Ridge Regression Equation is roughly:

$$E_{ridge} = \sum_{j=1}^n (-7.818x_{0j} + 9.918x_{1j} - y_j)^2 + 0.1((-7.818)^2 + 9.918^2)$$

Training Error for Ridge Regression: 1381.5759089148248

Test Error for Ridge Regression 280.1782039354006

Lasso Regression

Linear Lasso Regression Equatin is roughly:

$$E_{lasso} = \sum_{j=1}^n (-7.818x_{0j} + 9.918x_{1j} - y_j)^2 + 0.1(|-7.818| + |9.918|)$$

Training Error for Lasso Regression: 207.50080013508762

Test Error for Lasso Regression 266.0036661747493

Part f. I did obtain $\lambda = 0.1$ analytically. I used the code to compute

$$\lambda^* = \operatorname{argmin}_{\lambda} \left(\frac{1}{k} \sum_{i=1}^k \operatorname{loss}(g^{(D_k)}, y) \right)$$

```
[11]: final_errors_ridge = np.
      →array([lam_ridge[0][1],error_test_ridge,error_train_ridge])
final_errors_lasso = np.
      →array([lam_lasso[0][1],error_train_lasso,error_test_lasso])

final_errors_ridge_df = pd.DataFrame([final_errors_ridge],columns =_
      →['Cross-Validation Error','Testing Error','Training Error'],index = ['Ridge_
      →Regression'])
final_errors_lasso_df = pd.DataFrame([final_errors_lasso],columns =_
      →['Cross-Validation Error','Training Error','Testing Error'],index = ['Lasso_
      →Regression'])
print("Part g.\n")
print(final_errors_ridge_df,"\n\n",final_errors_lasso_df)
```

Part g.

	Cross-Validation Error	Testing Error	Training Error
Ridge Regression	60.763396	280.178204	1381.575909

	Cross-Validation Error	Training Error	Testing Error
Lasso Regression	40.779563	207.5008	266.003666

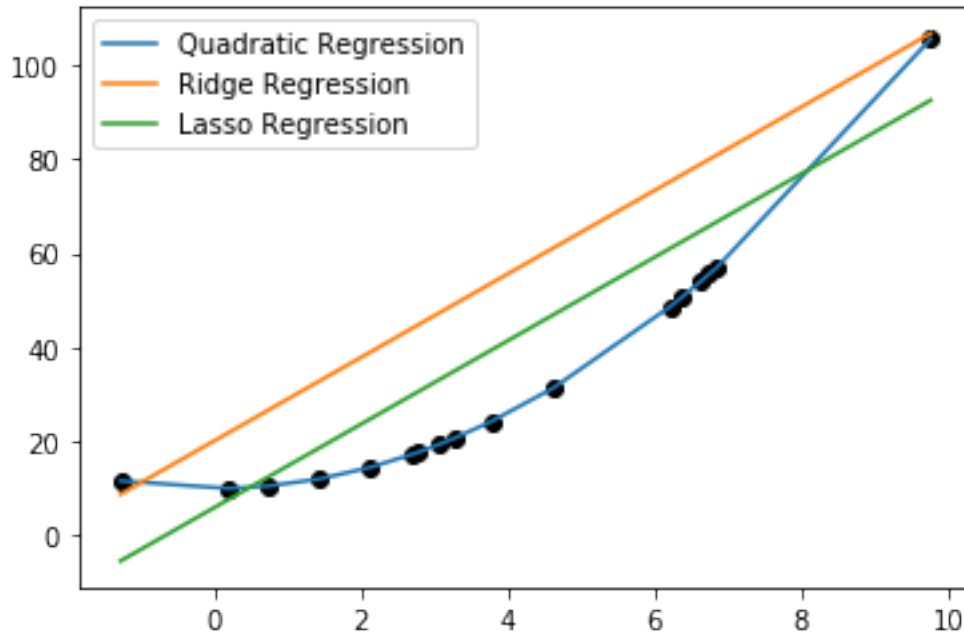
```
[12]: import matplotlib.pyplot as plt

def ridge_plot(x):
    return 4.22+8.85*x+15.949
def lasso_plot(x):
    return 4.22+8.85*x+1.774

x_vals = np.concatenate((x_train[:,1],x_test[:,1]))
y_vals = np.concatenate((y_train,y_test))
plt.scatter(x_vals,y_vals,color = "black")
x_vals = np.sort(x_vals)
target, = plt.plot(x_vals,target(x_vals))
ridge, = plt.plot(x_vals,ridge_plot(x_vals))
lasso, = plt.plot(x_vals,lasso_plot(x_vals))

plt.legend((target,ridge,lasso),('Quadratic Regression','Ridge_
Regression','Lasso Regression'))
```

[12]: <matplotlib.legend.Legend at 0x11465cc90>



[]: